

Final Project –CIS 668/IST 664

Classification of Kaggle Movie Reviews



Student:

Midhuneshwar Kandasamy Kanivalavan

SUID: 351564889

Mail: mkandasa@syr.edu

Professor: Michael Larche

Table of Contents

Introduction	4
Data Description	4
Reading the Train Data.....	6
Tokenization and Filtering	6
a. Pre processing documents	7
b. Retrieving words/tokens.....	8
Creating Feature Sets	8
a. Bag of Words (BOW)/Unigram features	9
b. Bigram Features	9
c. Sentiment Lexicons	9
d. Negation word features:	10
e. POS features.....	11
Classification: Naïve Bayes Classifier.....	11
a. Normal Features without preprocessing	12
b. Bigram Features	13
c. Negation word Features	14
d. Preprocessed Features.....	15
e. Sentiment Lexicon features	16
Combination of features sets.....	17
Comparative Analysis of Logistic Regression and Decision Tree classifier	18
Logistic Regression	18
a. Normal Features without preprocessing	20
b. Bigram Features	21
c. Negation Word Features.....	22
d. Preprocessed Features.....	23
e. Sentiment Lexicon Features.....	24
Decision Tree Classifier	25
a. Normal Features without preprocessing	26
b. Bigram Features	27
c. Negation Word Features.....	28
d. Preprocessed Features.....	29
e. Sentiment Lexicon Features.....	30

Highlights of Sci-kit learn algorithms	31
Comparison of evaluation metrics	31
Conclusion	31
Lessons Learned.....	32

Introduction

The primary objective of this project is to classify and predict the sentiment of movie reviews. The scale of sentiment classification is 0-4 with equal intervals. The details of scale for sentiment follows,

Classification Scale	Sentiment
0	Negative
1	Strong Negative
2	Neutral
3	Positive
4	Strong Positive

Based on the classification intervals, it can be said that the scaling is biased towards negative and positive reviews leaving neutral sentiment scoring a single interval classification.

The dataset was produced for the Kaggle competition, described here [data](#). This dataset data from the sentiment analysis by Socher et al, detailed at this web site: <http://nlp.stanford.edu/sentiment/>.

The data was taken from the original Pang and Lee movie review corpus based on reviews from the Rotten Tomatoes web site. Socher's group used crowd-sourcing to manually annotate all the subphrases of sentences with a sentiment label ranging over: "negative", "somewhat negative", "neutral", "somewhat positive", "positive".

Data Description


The data provided for this project has two data files namely 'train.csv' and 'test.csv'. The details of train data are,

- 1,56,060 rows of movie reviews. Each row represents review for a movie
- 4 columns, these are

Column_name	Description
Phraseld	Represents id number of a phrase
Sentenceld	Represents id for each sentence

Column_name	Description
Phrase	Represents the review for movie
Phrase Sentiment	Represents the sentiment for each review

Histogram of count of sentiments for train data,



The details of test data are ,

- 66,292 rows of movie reviews. Each row represents review for a movie
- 3 columns, these are

Column_name	Description
Phraseid	Represents id number of a phrase
Sentenceid	Represents id for each sentence
Phrase	Represents the review for movie

It must be noted that test data has no sentiment, which means that sentiment column is the outcome/target label which must be predicted. Therefore, for all of our analysis Phraseid, SentenceID and Phrase will be features and Sentiment will be target label.

Reading the Train Data

The first step is to read the data into python environment. This is done using the code below,

```
# function to read kaggle training file, train and test a classifier
def processkaggle(dirPath, limitStr):
    # convert the limit argument from a string to an int
    limit = int(limitStr)

    os.chdir(dirPath)

    f = open('./train.tsv', 'r')
    # loop over lines in the file and use the first limit of them
    phrasedata = []
    for line in f:
        # ignore the first line starting with Phrase and read all lines
        if (not line.startswith('Phrase')):
            # remove final end of line character
            line = line.strip()
            # each line has 4 items separated by tabs
            # ignore the phrase and sentence ids, and keep the phrase and sentiment
            phrasedata.append(line.split('\t')[2:4])
```

Here, we are writing a function which will loop through each line of the csv line and create a list called 'phrasedata'. Further, we have a option to limit the number of rows we can include in the loaded data. This limiting the loaded data is helpful when we want to run the analysis on a limited number of the instances. Since we are trying to classify and then predict, the better approach would be to load all the available data, however, the constraint is the computing power. Though, we have a constraint we tried to run our analysis on entire dataset.

Tokenization and Filtering

Once we have loaded the entire train dataset, we need to tokenize each word of the a phrase to proceed with our analysis. Using these tokenized phrases, we can create multiple different feature sets. The code developed is below,

```

run_sklearn_model_performance.py x run_sklearn_model_performance_Decision_Tree.py x modified_classifyKaggle.py x
313 # pick a random sample of length limit because of phrase overlapping sequences
314 random.shuffle(phrasedata)
315 phraselist = phrasedata[:limit]
316
317 print('Read', len(phrasedata), 'phrases, using', len(phraselist), 'random phrases')
318 #for phrase in phraselist[:10]:
319     #print (phrase)
320
321 # create list of phrase documents as (list of words, label)
322 phrasedocs = []
323 phrasedocs_without = []
324 # add all the phrases
325 for phrase in phraselist:
326
327     #without preprocessing
328     tokens = nltk.word_tokenize(phrase[0])
329     phrasedocs_without.append((tokens, int(phrase[1])))
330
331     # with pre processing
332     tokenizer = RegexpTokenizer(r'\w+')
333     phrase[0] = pre_processing_documents(phrase[0])
334     tokens = tokenizer.tokenize(phrase[0])
335     phrasedocs.append((tokens, int(phrase[1])))
336
337 # possibly filter tokens
338 normaltokens = get_words_from_phrasedocs_normal(phrasedocs_without)
339 preprocessedTokens = get_words_from_phrasedocs(phrasedocs)
340
341 # continue as usual to get all words and create word features
342 word_features = get_word_features(normaltokens)
343 featuresets_without_preprocessing = [(normal_features(d, word_features), s) for (d, s) in phrasedocs_without]
processkaggle()

```

From the code it can be observed that we are leveraging tokenize function of nltk to tokenize the phrases. Before tokenization, we are preprocessing for all those tokens using a regular expression '\w+', this regex pattern matches any alphanumeric with underscore with one or more occurrences. Once we have tokenized the data, we create two different lists 'normaltokens' which includes word phrases without filtering and the second list is 'preprocessedTokens' which includes word phrases with filtering.

To perform the above said tokenization and filtering we have defined the following functions,

a. Pre processing documents

This function splits each phrase into individual lines and converts them into lower case. Later, this lower case lines are run through regular expressions wherein any word with punctuation marks is removed and stored in a list called 'word_list'. Further this list is checked for any stop words i.e. all the stop words are removed from the list and a list called 'final_word_list' is created.

```

def pre_processing_documents(document):
    # "Pre_processing_documents"
    # "create list of lower case words"
    word_list = re.split('\s+', document.lower())
    # punctuation and numbers to be removed
    punctuation = re.compile(r'[-.?!\/%@,":;()|0-9]')
    word_list = [punctuation.sub("", word) for word in word_list]
    final_word_list = []
    for word in word_list:
        if word not in newstopwords:
            final_word_list.append(word)
    line = " ".join(final_word_list)
    return line

```

b. Retrieving words/tokens

There are three functions defined in this code. The first function returns list which contains tokens/words from the documents where the length is greater than 3. The second function returns a list of all words with corresponding sentiment. The third function returns all lines from the tokens.

```

def get_words_from_phasedocs(docs):
    all_words = []
    for (words, sentiment) in docs:
        # more than 3 length
        possible_words = [x for x in words if len(x) >= 3]
        all_words.extend(possible_words)
    return all_words

def get_words_from_phasedocs_normal(docs):
    all_words = []
    for (words, sentiment) in docs:
        all_words.extend(words)
    return all_words

# get all words from tokens
def get_words_from_test(lines):
    all_words = []
    for id, words in lines:
        all_words.extend(words)
    return all_words

```

Creating Feature Sets

The next key step in classification tasks in NLP is to create features from raw tokens. We had to define multiple features generating functions such as bag of words(BOW), bigrams, sentiment lexicons, negation words, POS tag features etc. Let us try to understand each of these functions,

a. Bag of Words (BOW)/Unigram features

For creating unigrams we defined two different functions. The first function creates a list of most repeated 200 words from the 'wordlist' which has processed tokens. The second function creates unique list of words from the documents and returns them as features.

```
def get_word_features(wordlist):
    wordlist = nltk.FreqDist(wordlist)
    word_features = [w for (w, c) in wordlist.most_common(200)]
    return word_features

def normal_features(document, word_features):
    document_words = set(document)
    features = {}
    for word in word_features:
        features['contains({})'.format(word)] = (word in document_words)
    return features
```

b. Bigram Features

We have worked on generating bigram feature from documents to get high frequent bigrams. We have filtered out special characters as well as filter by frequency. We have used the nbest function which just returns the highest scoring bigrams, using the number specified in both the measures.

```
def bigram_document_features(document, word_features, bigram_features):
    document_words = set(document)
    document_bigrams = nltk.bigrams(document)
    features = {}
    for word in word_features:
        features['contains({})'.format(word)] = (word in document_words)
    for bigram in bigram_features:
        features['bigram({} {})'.format(bigram[0], bigram[1])] = (bigram in document_bigrams)
    return features

def get_bigram_features(tokens):
    bigram_measures = nltk.collocations.BigramAssocMeasures()
    finder = BigramCollocationFinder.from_words(tokens, window_size=3)
    #finder.apply_freq_filter(6)
    bigram_features = finder.nbest(bigram_measures.chi_sq, 3000)
    return bigram_features[:500]
```

c. Sentiment Lexicons

We will first read in the subjectivity words from the subjectivity lexicon file created by Janyce Wiebe and her group at the University of Pittsburgh in the MPQA project. Although these words are often

used as features themselves or in conjunction with other information, we will create two features that involve counting the positive and negative subjectivity words present in each document. I copy and pasted the definition of the readSubjectivity function from the Subjectivity.py module which is provided by Professor. It creates a Subjectivity Lexicon that is represented here as a dictionary, where each word is mapped to a list containing the strength and polarity. A feature extraction function that has all the word features as before, but also has two features 'positivecount' and 'negativecount'. These features contain counts of all the positive and negative subjectivity words, where each weakly subjective word is counted once and each strongly subjective word is counted twice.

d. Negation word features:

Negation of opinions is an important part of sentimental classification. Here I tried a simple strategy which professor explained in Lab-10. I look for negation words "not", "never" and "no" and negation that appears in contractions of the form "doesn", "", "t". For example, my first document has the following words: if, 'you', 'don', "", 't', 'like', 'this', 'film', ',', 'then', 'you', 'have', 'a', 'problem', 'with', 'the', 'genre', 'itself'. One strategy with negation words is to negate the word following the negation word, while other strategies negate all words up to the next punctuation or use syntax to find the scope of the negation. I followed the first strategy here, and I go through the document words in order adding the word features, but if the word follows a negation words, change the feature to negated word.

```
negationwords = ['no', 'not', 'never', 'none', 'nowhere', 'nothing', 'noone', 'rather',
                 'hardly', 'scarcely', 'rarely', 'seldom', 'neither', 'nor']
def NOT_features(document, word_features, negationwords):
    features = {}
    for word in word_features:
        features['contains({})'.format(word)] = False
        features['contains(NOT{})'.format(word)] = False
    # go through document words in order
    for i in range(0, len(document)):
        word = document[i]
        if ((i + 1) < len(document)) and (word in negationwords):
            i += 1
            features['contains(NOT{})'.format(document[i])] = (document[i] in word_features)
        else:
            if ((i + 3) < len(document)) and (word.endswith('\n') and document[i+1] == "" and document[i+2] == 't'):
                i += 3
                features['contains(NOT{})'.format(document[i])] = (document[i] in word_features)
            else:
                features['contains({})'.format(word)] = (word in word_features)
    return features
```

e. POS features

We have done this classification task with help of part-of-speech tag features. This is more likely for shorter units of classification; such as sentence level classification or shorter social media such as tweets. In this dataset, we have large training dataset and moreover, in the NLTK, this is difficult to demonstrate, since on computer, it takes the default NLTK POS tagger too much time. Because of this limitation we tested on only 2000 training sentences. The most common way to use POS tagging information is to include counts of various types of word tags.

```
def POS_features(document, word_features):
    document_words = set(document)
    tagged_words = nltk.pos_tag(document)
    features = {}
    for word in word_features:
        features['contains({})'.format(word)] = (word in document_words)
    numNoun = 0
    numVerb = 0
    numAdj = 0
    numAdverb = 0
    for (word, tag) in tagged_words:
        if tag.startswith('N'): numNoun += 1
        if tag.startswith('V'): numVerb += 1
        if tag.startswith('J'): numAdj += 1
        if tag.startswith('R'): numAdverb += 1
    features['nouns'] = numNoun
    features['verbs'] = numVerb
    features['adjectives'] = numAdj
    features['adverbs'] = numAdverb
    return features
```

Classification: Naïve Bayes Classifier

In this section we will discuss the classifier obtained for each of the feature sets created by running the above functions. Before delving into details of outputs for each of the feature sets, let's look at the accuracies obtained for each of them,

Feature Set	Accuracy achieved
Normal Features without preprocessing	51.86%
Bigram Features	53.75%
Negation word Features	55.47%
Preprocessed Features	53.75%
Sentiment Lexicon features	55.08%

Since, we have established the accuracies for each of the feature sets, let us now focus on the confusion matrix and make some definitive conclusions.

a. Normal Features without preprocessing

```

ice.py x run_sklern_model_performance_Decision_Tree.py x modified_classifyKaggle.py x Local x
337 normaltokens = get_words_from_phraseset(normaltokens)
340 preprocessedTokens = get_words_from_phraseset(preprocessedTokens)
341
342 # continue as usual to get all words and create word features
343
344
345
346 word_features = get_word_features(normaltokens)
347 featuresets_without_preprocessing = [(normal_features(d, word_features), s) for (d, s) in corpus]
348 #print featuresets_without_preprocessing[0]
349 writeFeatureSets(featuresets_without_preprocessing, "features_without_preprocessing.csv")
350 print ("-----")
351 print ("Accuracy with normal features, without pre-processing steps : ")
352 accuracy_calculation(featuresets_without_preprocessing)
353
354
355 word_features = get_word_features(preprocessedTokens)
356 #print word_features[:20]
357 featuresets = [(normal_features(d, word_features), s) for (d, s) in corpus]
358 #print featuresets[0]
359 writeFeatureSets(featuresets, "features_preprocessed.csv")
360 print ("-----")
361 print ("Accuracy with pre-processed features : ")
362 accuracy_calculation(featuresets)
363
364
365 SL_featuresets = [(SL_features(d, word_features, SL), c) for (d, s, c) in corpus]
366 writeFeatureSets(SL_featuresets, "features_SL.csv")
367 #print SL_featuresets[0]
368 print ("-----")
369 print ("Accuracy with SL_featuresets : ")
370 accuracy_calculation(SL_featuresets)
371
372
373 processkaggle() for phrase in phraseset

```

(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemovieireviews>python modified_classifyKaggle.py "C:\Users\pkous\Desktop\WorkingDir NLP\kagglemovieireviews\corpus" 156000

Read 156060 phrases, using 156000 random phrases

Accuracy with normal features, without pre-processing steps :
Training and testing a classifier
Accuracy of classifier : 0.5186538461538461

Showing most informative features

Most Informative Features

contains(bad) = True	0 : 4	=	82.7 : 1.0
contains(minutes) = True	0 : 4	=	22.9 : 1.0
contains(best) = True	4 : 2	=	16.8 : 1.0
contains(entertaining) = True	4 : 2	=	16.0 : 1.0
contains(performances) = True	4 : 2	=	15.2 : 1.0
contains(great) = True	4 : 2	=	14.0 : 1.0
contains(fun) = True	4 : 2	=	12.6 : 1.0
contains(Love) = True	4 : 0	=	12.2 : 1.0
contains(funny) = True	4 : 2	=	11.5 : 1.0
contains(performance) = True	4 : 2	=	11.2 : 1.0
contains(too) = True	1 : 4	=	10.4 : 1.0
contains(plot) = True	0 : 4	=	9.3 : 1.0
contains(A) = True	4 : 2	=	7.8 : 1.0
contains(heart) = True	4 : 1	=	7.3 : 1.0
contains(acting) = True	0 : 2	=	7.0 : 1.0
contains(humor) = True	4 : 2	=	6.6 : 1.0
contains(would) = True	0 : 4	=	6.5 : 1.0
contains(well) = True	4 : 2	=	6.1 : 1.0
contains(most) = True	4 : 2	=	6.1 : 1.0
contains(nothing) = True	0 : 2	=	5.9 : 1.0
contains(script) = True	0 : 2	=	5.9 : 1.0

Python Console
interpreter // Configure a Python interpreter... (today 7:48 PM)

327:1 LF UTF-8 2 spaces* Python 3.8 (base)

The confusion matrix						Classification	Sentiment
	0	1	2	3	4	Scale	
0	121	319	87	35		0	Negative
1	310	1523	266	75		1	Strong Negative
2	177	501	6648	541	81	2	Neutral
3	212	288	2031	606	204	3	Positive
4	65	52	403	218	147	4	Strong Positive
(row = reference; col = test)							

From the above confusion matrix, it can be said that most of neutral sentiments are predicted accurately and the least accurately predicted label is positive sentiment labels.

b. Bigram Features

```

378 bigram_features = get_bigram_features(preprocessedTokens)
379 #print bigram_features[0]
380 bigram_featuresets = [(bigram_document_features(d, word_featuresets[0]), c) for d, c in bigram_features]
381 #print bigram_featuresets[0]
382 writeFeatureSets(bigram_featuresets, "features_biagram.csv")
383 print ("-----")
384 print ("Accuracy with bigram featuresets : ")
385 accuracy_calculation(bigram_featuresets)
386
387
388
389 '''pos_featuresets = [(POS_features(d, word_featuresets[0]), c) for d, c in pos_features]
390 print pos_featuresets[0]
391 writeFeatureSets(pos_featuresets, "features_pos.csv")
392 print "-----"
393 print "Accuracy with pos_featuresets : "
394 writeFeatureSets(pos_featuresets, "features_pos_featuresets.csv")
395 accuracy_calculation(pos_featuresets)'''
396
397
398
399 #####
400 ##### generate test csv file. #####
401 #####
402
403 f = open('./test.tsv', 'r')
404 # loop over lines in the file and use the first limit of them
405 testphrasedata = []
406 for line in f:
407     # ignore the first line starting with Phrase and read all lines
408     if (not line.startswith('Phrase')):
409         # remove final end of line character
410         line = line.rstrip()
411         testphrasedata.append(line)
412
413 processkaggle()

```

```

-----
Accuracy with bigram featuresets :
Training and testing a classifier
Accuracy of classifier :
0.5375641025641026
-----
Showing most informative features
Most Informative Features
contains(bad) = True          0 : 4      = 72.0 : 1.0
contains(moving) = True       4 : 0      = 36.7 : 1.0
contains(dull) = True         0 : 2      = 31.8 : 1.0
contains(fascinating) = True  4 : 2      = 21.5 : 1.0
contains(minutes) = True      0 : 4      = 18.6 : 1.0
contains(best) = True         4 : 2      = 16.2 : 1.0
contains(entertaining) = True 4 : 2      = 15.5 : 1.0
contains(performances) = True 4 : 2      = 14.9 : 1.0
contains(great) = True        4 : 2      = 13.7 : 1.0
contains(fun) = True          4 : 2      = 12.6 : 1.0
contains(sweet) = True        4 : 2      = 11.9 : 1.0
contains(funny) = True        4 : 2      = 11.6 : 1.0
contains(performance) = True  4 : 2      = 11.2 : 1.0
contains(dialogue) = True     0 : 2      = 10.7 : 1.0
contains(compelling) = True   4 : 2      = 10.5 : 1.0
contains(piece) = True        4 : 2      = 9.5 : 1.0
contains(plot) = True         0 : 4      = 9.3 : 1.0
contains(think) = True        0 : 4      = 8.7 : 1.0
contains(year) = True         4 : 2      = 8.7 : 1.0
contains(series) = True       0 : 4      = 7.6 : 1.0
contains(experience) = True   4 : 2      = 7.4 : 1.0
contains(heart) = True        4 : 0      = 7.3 : 1.0
contains(feature) = True      4 : 0      = 7.0 : 1.0

```

Python Console
interpreter // Configure a Python interpreter... (today 7:48 PM)

412:1 LF UTF-8 2 spaces* Python 3.8 (base)

The confusion matrix

	0	1	2	3	4
0	<46>	127	483	68	8
1	49	<338>	2038	241	28
2	25	283	<7199>	410	31
3	22	171	2336	<724>	88
4	6	42	486	272	<79>

(row = reference; col = test)

1 Event Log

412:1 LF UTF-8 2 spaces* Python 3.8 (base)

Classification	Sentiment
0	Negative
1	Strong Negative
2	Neutral
3	Positive
4	Strong Positive

From the above confusion matrix, it can be said that most of neutral sentiments and strong negative labels are predicted accurately for most of the training data and the least accurately predicted label is positive sentiment labels.

d. Preprocessed Features

```

355 word_features = get_word_features(preprocessed_tokens)
356 #print word_features[:20]
357 featuresets = [(normal_features(d, word_features), s) for (d, s) in data]
358 #print featuresets[0]
359 writeFeatureSets(featuresets, "features_preprocessed.csv")
360 print ("-----")
361 print ("Accuracy with pre-processed features : ")
362 accuracy_calculation(featuresets)
363
364
365 SL_featuresets = [(SL_features(d, word_features, SL), c) for (d, s, c) in data]
366 writeFeatureSets(SL_featuresets, "features_SL.csv")
367 #print SL_featuresets[0]
368 print ("-----")
369 print ("Accuracy with SL_featuresets : ")
370 accuracy_calculation(SL_featuresets)
371
372 NOT_featuresets = [(NOT_features(d, word_features, negationwords), c) for (d, s, c) in data]
373 #print NOT_featuresets[0]
374 writeFeatureSets(SL_featuresets, "features_NOT.csv")
375 print ("-----")
376 print ("Accuracy with NOT_featuresets : ")
377 accuracy_calculation(NOT_featuresets)
378
379 bigram_features = get_bigram_features(preprocessed_tokens)
380 #print bigram_features[0]
381 bigram_featuresets = [(bigram_document_features(d, word_features, bigram_features), s) for (d, s) in data]
382 #print bigram_featuresets[0]
383 writeFeatureSets(bigram_featuresets, "features_bigram.csv")
384 print ("-----")
385 print ("Accuracy with bigram featuresets : ")
386 accuracy_calculation(bigram_featuresets)

```

Accuracy with pre-processed features :
Training and testing a classifier
Accuracy of classifier :
0.5375641025641026

Showing most informative features

Most Informative Features

contains(bad) = True	0 : 4	=	72.0 : 1.0
contains(moving) = True	4 : 0	=	36.7 : 1.0
contains(dull) = True	0 : 2	=	31.8 : 1.0
contains(fascinating) = True	4 : 2	=	21.5 : 1.0
contains(minutes) = True	0 : 4	=	18.6 : 1.0
contains(best) = True	4 : 2	=	16.2 : 1.0
contains(entertaining) = True	4 : 2	=	15.5 : 1.0
contains(performance) = True	4 : 2	=	14.9 : 1.0
contains(great) = True	4 : 2	=	13.7 : 1.0
contains(fun) = True	4 : 2	=	12.6 : 1.0
contains(sweet) = True	4 : 2	=	11.9 : 1.0
contains(funny) = True	4 : 2	=	11.6 : 1.0
contains(performance) = True	4 : 2	=	11.2 : 1.0
contains(dialogue) = True	0 : 2	=	10.7 : 1.0
contains(compelling) = True	4 : 2	=	10.5 : 1.0
contains(piece) = True	4 : 2	=	9.5 : 1.0
contains(plot) = True	0 : 4	=	9.3 : 1.0
contains(think) = True	0 : 4	=	8.7 : 1.0
contains(year) = True	4 : 2	=	8.7 : 1.0
contains(series) = True	0 : 4	=	7.6 : 1.0
contains(experience) = True	4 : 2	=	7.4 : 1.0
contains(heart) = True	4 : 0	=	7.3 : 1.0
contains(feature) = True	4 : 0	=	7.0 : 1.0

The confusion matrix

	0	1	2	3	4
0	<46> 127	483	68	8	
1	49	<338>2038	241	28	
2	25	283	<7199> 410	31	
3	22	171	2336	<724> 88	
4	6	42	486	272	<79>

(row = reference; col = test)

Classification Scale	Sentiment
0	Negative
1	Strong Negative
2	Neutral
3	Positive
4	Strong Positive

From the above confusion matrix, neutral sentiment labels are predicted accurately.

e. Sentiment Lexicon features

```

363
364
365 SL_featuresets = [(SL_features(d, word_features, SL), c) for (d, c) in dataset]
366 writeFeatureSets(SL_featuresets, "features_SL.csv")
367 #print SL_featuresets[0]
368 print ("-----")
369 print ("Accuracy with SL_featuresets : ")
370 accuracy_calculation(SL_featuresets)
371
372 NOT_featuresets = [(NOT_features(d, word_features, negationword), c) for (d, c) in dataset]
373 #print NOT_featuresets[0]
374 writeFeatureSets(SL_featuresets, "features_NOT.csv")
375 print ("-----")
376 print ("Accuracy with NOT_featuresets : ")
377 accuracy_calculation(NOT_featuresets)
378
379 bigram_features = get_bigram_features(preprocessedTokens)
380 #print bigram_features[0]
381 bigram_featuresets = [(bigram_document_features(d, word_features, bigram_features), c) for (d, c) in dataset]
382 #print bigram_featuresets[0]
383 writeFeatureSets(bigram_featuresets, "features_bigram.csv")
384 print ("-----")
385 print ("Accuracy with bigram_featuresets : ")
386 accuracy_calculation(bigram_featuresets)
387
388
389 pos_featuresets = [(POS_features(d, word_features, c) for (d, c) in dataset]
390 print pos_featuresets[0]
391 writeFeatureSets(pos_featuresets, "features_pos.csv")
392 print "-----"
393 print "Accuracy with pos_featuresets : "
394 writeFeatureSets(pos_featuresets, "features_pos_featuresets.csv")
395
processkaggle() for phrase in phraselist

```

```

Accuracy with SL_featuresets :
Training and testing a classifier
Accuracy of classifier :
0.5508974358974359

Showing most informative features
Most Informative Features
contains(bad) = True          0 : 4      = 72.0 : 1.0
positivecount = 11           4 : 2      = 71.7 : 1.0
positivecount = 10           4 : 2      = 44.9 : 1.0
negativecount = 11           0 : 2      = 41.4 : 1.0
contains(moving) = True       4 : 0      = 36.7 : 1.0
contains(dull) = True         0 : 2      = 31.8 : 1.0
negativecount = 9             0 : 2      = 29.6 : 1.0
positivecount = 9             4 : 1      = 28.7 : 1.0
negativecount = 8             0 : 2      = 27.3 : 1.0
contains(fascinating) = True  4 : 2      = 21.5 : 1.0
contains(minutes) = True     0 : 4      = 18.6 : 1.0
negativecount = 10           0 : 2      = 17.9 : 1.0
positivecount = 12           4 : 2      = 17.8 : 1.0
contains(best) = True        4 : 2      = 16.2 : 1.0
positivecount = 8             4 : 2      = 16.1 : 1.0
positivecount = 6             4 : 2      = 15.7 : 1.0
positivecount = 7             4 : 2      = 15.5 : 1.0
contains(entertaining) = True 4 : 2      = 15.5 : 1.0
contains(performances) = True 4 : 2      = 14.9 : 1.0
negativecount = 7             0 : 2      = 14.8 : 1.0
contains(great) = True       4 : 2      = 13.7 : 1.0
contains(fun) = True         4 : 2      = 12.6 : 1.0
positivecount = 5             4 : 2      = 12.6 : 1.0
contains(sweet) = True       4 : 2      = 11.9 : 1.0
contains(funny) = True       4 : 2      = 11.6 : 1.0

```

Python Console
interpreter // Configure a Python interpreter... (today 7:48 PM)

<p>The confusion matrix</p> <pre> 0 1 2 3 4 +-----+ 0 <90> 162 370 89 21 1 108 <478>1692 357 59 2 50 389<6732> 719 58 3 29 196 1781<1138> 197 4 10 40 288 391 <156> +-----+ (row = reference; col = test) </pre>	<table> <tr> <th>Classification Scale</th><th>Sentiment</th></tr> <tr> <td>0</td><td>Negative</td></tr> <tr> <td>1</td><td>Strong Negative</td></tr> <tr> <td>2</td><td>Neutral</td></tr> <tr> <td>3</td><td>Positive</td></tr> <tr> <td>4</td><td>Strong Positive</td></tr> </table>	Classification Scale	Sentiment	0	Negative	1	Strong Negative	2	Neutral	3	Positive	4	Strong Positive
Classification Scale	Sentiment												
0	Negative												
1	Strong Negative												
2	Neutral												
3	Positive												
4	Strong Positive												

From the above confusion matrix, neutral and positive sentiment labels are accurately predicted.

Combination of features sets:

In this section, we will try to create a new function which will combine different feature sets such as Sentiment Lexicons, Bigram features and unigram features. The code follows,

```
def combined_document_features(document, word_features, SL, bigram_features):
    document_words = set(document)
    document_bigrams = nltk.bigrams(document)
    features = {}

    for word in document_words:
        # features object
        posword = 0
        neutword = 0
        negword = 0
        for word in document_words:
            if word in SL[0]:
                posword += 1
            if word in SL[1]:
                neutword += 1
            if word in SL[2]:
                negword += 1
            features['positivecount'] = posword
            features['neutralcount'] = neutword
            features['negativecount'] = negword

        for word in word_features:
            features['V_{}'.format(word)] = False
            features['V_NOT{}'.format(word)] = False

        for bigram in bigram_features:
            features['B_{}_{}'.format(bigram[0], bigram[1])] = (bigram in document_bigrams)

    return features
```

From the above code we tried to create one single feature set which is a combination of bigrams and sentiment lexicons.

```
Accuracy with combined featuresets :
Training and testing a classifier
Accuracy of classifier :
0.8
```

The accuracy we obtained for combined features is 80%. One of possible reasons for such high accuracy is because we are capturing the non-linearity in the data by creating and combining multiple features because of which the algorithm is maximizing learning.

The main constraint running this function is that the processing time to generate feature sets was time taking i.e. the execution time was high.

Comparative Analysis of Logistic Regression and Decision Tree classifier:

In this section we will try to use Sci-kit learn algorithms to classify the sentiments. Since our outcome labels are like classification type of labels, we have decided to implement Logistic Regression and Decision tree classifier algorithms. Further we will run both the algorithms on each of the five feature sets.

Before delving into any details, we will do a comparative analysis of the metrics,

Feature set type	Logistic Regression			Decision Tree Classifier		
	Precision	Recall	f1-score	Precision	Recall	f1-score
Normal Features without preprocessing	0.49	0.47	0.46	0.47	0.53	0.42
Preprocessed Features	0.46	0.44	0.43	0.49	0.52	0.38
Bigram Features	0.46	0.44	0.43	0.49	0.52	0.38
Negation word Features	0.53	0.51	0.51	0.51	0.55	0.50
Sentiment Lexicon features	0.53	0.50	0.51	0.51	0.55	0.50

For Logistic regression classifier, Sentiment Lexicon features performance better in classification compared with other features functions because fewer words are unseen in train data as features. These words or tokens covers on Lexicon dictionary. Plus, we observed recall score lesser than F-measure which is greater than Precision.

For Decision Tree classifier, Sentiment Lexicon feature and Negation word features perform better in classification compared with other features functions. These words or tokens covers on Lexicon dictionary. Plus, we observed recall score higher than F-measure which is lesser than Precision.

Let us now discuss outputs for each of feature set for an algorithm,

Logistic Regression:

The parameters used in logistic regression are,

- **Class_weight:** Weights associated with classes in the form {class_label: weight}. If not given, all classes are supposed to have weight one. The “balanced” mode uses the values of y to automatically adjust weights inversely proportional to class frequencies in the input data as $n_{\text{samples}} / (n_{\text{classes}} * \text{np.bincount}(y))$.
- **solver:** Algorithm to use in the optimization problem.

- max_iter: Maximum number of iterations taken for the solvers to converge.

For our analysis, we have used the following parameters,

Parameter_Description	Parameter_value
Class_weight	Balanced
Solver	Lbfgs
max_iter	1000

We have done experiments on different feature sets created using Logistic regression algorithm,

a. Normal Features without preprocessing

```
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>python run_sklearn_model_performance.py "C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews\corpus\features_normal.csv"
```

```
Shape of feature data - num instances with num features + class label  
(156000, 201)
```

```
** Results from Logistic Regression with liblinear
```

	precision	recall	f1-score	support
neg	0.15	0.39	0.22	7069
neu	0.66	0.68	0.67	79548
pos	0.20	0.48	0.28	9202
sneg	0.33	0.20	0.25	27263
spos	0.34	0.17	0.23	32918
accuracy			0.47	156000
macro avg	0.34	0.39	0.33	156000
weighted avg	0.49	0.47	0.46	156000

Predicted	neg	neu	pos	sneg	spos	All
Actual						
neg	2761	1681	1165	991	471	7069
neu	5601	54437	5607	6942	6961	79548
pos	1087	2097	4398	492	1128	9202
sneg	5437	10812	2979	5531	2504	27263
spos	3370	13470	7702	2667	5709	32918
All	18256	82497	21851	16623	16773	156000

```
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>
```

Event Log

41:39 LF UTF-8 4 spaces Python 3.8 (base)

b. Bigram Features

```
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>python run_sklearn_model_performance.py "C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews\corpus\features_biagram.csv"
```

```
Shape of feature data - num instances with num features + class label  
(156000, 701)
```

```
** Results from Logistic Regression with liblinear
```

	precision	recall	f1-score	support
neg	0.13	0.37	0.20	7069
neu	0.62	0.67	0.64	79548
pos	0.19	0.45	0.27	9202
sneg	0.28	0.16	0.20	27263
spos	0.35	0.15	0.21	32918
accuracy			0.44	156000
macro avg	0.31	0.36	0.30	156000
weighted avg	0.46	0.44	0.43	156000

Predicted \ Actual	neg	neu	pos	sneg	spos	All
neg	2628	2418	845	847	331	7069
neu	7356	52937	6702	6888	5665	79548
pos	934	2574	4162	572	960	9202
sneg	5376	12723	2965	4325	1874	27263
spos	3354	14559	7461	2769	4775	32918
All	19648	85211	22135	15401	13605	156000

```
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>
```

1 Event Log

41:39 LF UTF-8 4 spaces Python 3.8 (base)

c. Negation Word Features

```

Terminal: Local x +
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>python run_sklearn_model_performance.py "C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews\corpus\features_NOT.csv"
Shape of feature data - num instances with num features + class label
(156000, 203)
** Results from Logistic Regression with liblinear
      precision    recall  f1-score   support

     neg         0.18         0.47         0.26         7069
     neu         0.71         0.67         0.69        79548
     pos         0.25         0.55         0.35         9202
    sneg         0.34         0.26         0.30        27263
    spos         0.41         0.27         0.33        32918

 accuracy                   0.50        156000
 macro avg         0.38         0.45         0.39        156000
 weighted avg         0.53         0.50         0.51        156000

Predicted   neg    neu    pos    sneg    spos    All
Actual
neg         3343   1247    505   1480    494    7069
neu        5453  53376   4038   9015   7666   79548
pos         575   1006   5058    410   2153    9202
sneg       6225   9372   2014   7121   2531   27263
spos       2615  10299   8258   2722   9024   32918
All       18211  75300  19873  20748  21868  156000

(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>

```

1 Event Log
41:39 LF UTF-8 4 spaces Python 3.8 (base)

d. Preprocessed Features

```
Terminal: Local x +
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>python run_sklearn_model_performance.py "C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews\corpus\features_preprocessed.csv"
Shape of feature data - num instances with num features + class label
(156000, 201)
** Results from Logistic Regression with liblinear
```

	precision	recall	f1-score	support
neg	0.13	0.37	0.20	7069
neu	0.62	0.67	0.64	79548
pos	0.19	0.45	0.27	9202
sneg	0.28	0.16	0.20	27263
spos	0.35	0.15	0.21	32918
accuracy			0.44	156000
macro avg	0.31	0.36	0.30	156000
weighted avg	0.46	0.44	0.43	156000

Predicted	neg	neu	pos	sneg	spos	All
Actual						
neg	2626	2418	846	847	332	7069
neu	7367	52924	6693	6887	5677	79548
pos	934	2574	4162	572	960	9202
sneg	5379	12723	2959	4327	1875	27263
spos	3359	14549	7459	2769	4782	32918
All	19665	85188	22119	15402	13626	156000

```
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>
```

1 Event Log

41:39 LF UTF-8 4 spaces Python 3.8 (base)

e. Sentiment Lexicon Features

```

Terminal: Local x +
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>python run_sklearn_model_performance.py "C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews\corpus\features_SL.csv"
Shape of feature data - num instances with num features + class label
(156000, 203)
** Results from Logistic Regression with liblinear
      precision    recall  f1-score   support

     neg         0.18         0.47         0.26         7069
     neu         0.71         0.67         0.69        79548
     pos         0.25         0.55         0.35         9202
    sneg         0.34         0.26         0.30        27263
    spos         0.41         0.27         0.33        32918

 accuracy                   0.50        156000
 macro avg         0.38         0.45         0.39        156000
 weighted avg         0.53         0.50         0.51        156000

Predicted   neg    neu    pos    sneg    spos    All
Actual
neg         3343   1247    505   1480    494    7069
neu        5453  53376   4038   9015   7666   79548
pos         575   1006   5058    410   2153    9202
sneg       6225   9372   2014   7121   2531   27263
spos       2615  10299   8258   2722   9024   32918
All       18211  75300  19873  20748  21868  156000

(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>

```

1 Event Log

262:10 (2 chars) LF UTF-8 2 spaces* Python 3.8 (base)

Decision Tree Classifier

The parameters used in Decision Tree Classifier are,

For our analysis, we have used the following parameters,

- **criterion:** The function to measure the quality of a split. Supported criteria are “gini” for the Gini impurity and “entropy” for the information gain.
- **max_depth:** The maximum depth of the tree. If None, then nodes are expanded until all leaves are pure or until all leaves contain less than min_samples_split samples.
- **min_samples_split:** The minimum number of samples required to split an internal node

Parameter_Description	Parameter_value
criterion	gini
max_depth	7
min_samples_split	5

We have done experiments on different feature sets created using Logistic regression algorithm,

a. Normal Features without preprocessing

```
Terminal: Local x +
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>python run_sklearn_
model_performance_Decision_Tree.py "C:\Users\pkous\Desktop\WorkingDir NLP\kagglemov
iereviews\corpus\features_normal.csv"
Shape of feature data - num instances with num features + class label
(156000, 201)
** Results from Decision Tree Classifier
```

	precision	recall	f1-score	support
neg	0.39	0.02	0.04	7069
neu	0.56	0.94	0.70	79548
pos	0.41	0.01	0.03	9202
sneg	0.43	0.05	0.10	27263
spos	0.32	0.17	0.22	32918
accuracy			0.53	156000
macro avg	0.42	0.24	0.22	156000
weighted avg	0.47	0.53	0.42	156000

Predicted	neg	neu	pos	sneg	spos	All
Actual						
neg	155	4745	7	457	1705	7069
neu	60	74620	34	903	3931	79548
pos	5	6115	120	95	2867	9202
sneg	157	21860	30	1465	3751	27263
spos	25	26657	104	453	5679	32918
All	402	133997	295	3373	17933	156000

```
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>
```

1 Event Log

7:1 LF UTF-8 4 spaces Python 3.8 (base)

b. Bigram Features

```

Terminal: Local x +
FileNotFoundError: [Errno 2] No such file or directory: 'C:\\Users\\pkous\\Desktop\\WorkingDir NLP\\kagglemoviereviews\\corpus\\features_bigram.csv'

(base) C:\\Users\\pkous\\Desktop\\WorkingDir NLP\\kagglemoviereviews>python run_sklearn_model_performance_Decision_Tree.py "C:\\Users\\pkous\\Desktop\\WorkingDir NLP\\kagglemoviereviews\\corpus\\features_biagram.csv"
Shape of feature data - num instances with num features + class label
(156000, 701)
** Results from Decision Tree Classifier

```

	precision	recall	f1-score	support
neg	0.55	0.02	0.03	7069
neu	0.52	0.99	0.68	79548
pos	0.51	0.03	0.05	9202
sneg	0.43	0.02	0.03	27263
spos	0.45	0.06	0.10	32918
accuracy			0.52	156000
macro avg	0.49	0.22	0.18	156000
weighted avg	0.49	0.52	0.38	156000

Predicted \ Actual	neg	neu	pos	sneg	spos	All
neg	108	6498	7	343	113	7069
neu	24	78375	61	237	851	79548
pos	1	8088	246	10	857	9202
sneg	58	26196	35	492	482	27263
spos	5	30821	134	73	1885	32918
All	196	149978	483	1155	4188	156000

```

(base) C:\\Users\\pkous\\Desktop\\WorkingDir NLP\\kagglemoviereviews>

```

1 Event Log

7:1 LF UTF-8 4 spaces Python 3.8 (base)

c. Negation Word Features

```
Terminal: Local x +
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>python run_sklearn_
model_performance_Decision_Tree.py "C:\Users\pkous\Desktop\WorkingDir NLP\kagglemov
iereviews\corpus\features_NOT.csv"
Shape of feature data - num instances with num features + class label
(156000, 203)
** Results from Decision Tree Classifier
```

	precision	recall	f1-score	support
neg	0.44	0.03	0.06	7069
neu	0.63	0.83	0.71	79548
pos	0.42	0.04	0.07	9202
sneg	0.36	0.16	0.23	27263
spos	0.41	0.46	0.43	32918
accuracy			0.55	156000
macro avg	0.45	0.30	0.30	156000
weighted avg	0.51	0.55	0.50	156000

Predicted \ Actual	neg	neu	pos	sneg	spos	All
neg	242	3415	21	2224	1167	7069
neu	67	66099	90	3551	9741	79548
pos	16	2201	334	325	6326	9202
sneg	176	18117	51	4452	4467	27263
spos	45	15799	297	1719	15058	32918
All	546	105631	793	12271	36759	156000

```
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>
```

7:1 LF UTF-8 4 spaces Python 3.8 (base)

d. Preprocessed Features

```

Terminal: Local x +
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>python run_sklearn_
model_performance_Decision_Tree.py "C:\Users\pkous\Desktop\WorkingDir NLP\kagglemov
iereviews\corpus\features_preprocessed.csv"
Shape of feature data - num instances with num features + class label
(156000, 201)
** Results from Decision Tree Classifier

```

	precision	recall	f1-score	support
neg	0.55	0.02	0.03	7069
neu	0.52	0.99	0.68	79548
pos	0.51	0.03	0.05	9202
sneg	0.43	0.02	0.03	27263
spos	0.45	0.06	0.10	32918
accuracy			0.52	156000
macro avg	0.49	0.22	0.18	156000
weighted avg	0.49	0.52	0.38	156000

```


```

Predicted	neg	neu	pos	sneg	spos	All
Actual						
neg	108	6498	7	343	113	7069
neu	24	78374	62	237	851	79548
pos	1	8088	246	10	857	9202
sneg	58	26196	33	494	482	27263
spos	5	30819	134	73	1887	32918
All	196	149975	482	1157	4190	156000

```

(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>

```

7:1 LF UTF-8 4 spaces Python 3.8 (base)

e. Sentiment Lexicon Features

```
Terminal: Local x +
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>python run_sklearn_model_performance_Decision_Tree.py "C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews\corpus\features_SL.csv"
Shape of feature data - num instances with num features + class label
(156000, 203)
** Results from Decision Tree Classifier
```

	precision	recall	f1-score	support
neg	0.44	0.03	0.06	7069
neu	0.63	0.83	0.71	79548
pos	0.42	0.04	0.07	9202
sneg	0.36	0.16	0.23	27263
spos	0.41	0.46	0.43	32918
accuracy			0.55	156000
macro avg	0.45	0.30	0.30	156000
weighted avg	0.51	0.55	0.50	156000

Predicted	neg	neu	pos	sneg	spos	All
Actual						
neg	242	3415	21	2224	1167	7069
neu	68	66099	90	3550	9741	79548
pos	16	2201	333	326	6326	9202
sneg	178	18118	51	4449	4467	27263
spos	45	15799	298	1719	15057	32918
All	549	105632	793	12268	36758	156000

```
(base) C:\Users\pkous\Desktop\WorkingDir NLP\kagglemoviereviews>
```

7:1 LF UTF-8 4 spaces Python 3.8 (base)

Highlights of Sci-kit learn algorithms:

After running through different experiments, it can be said that Sentiment Lexicon feature sets work best when classifying the movie review datasets. The important aspect is that both logistic regression and decision tree classifier work similarly for Sentiment Lexicon features with highest precision, recall and f1-measures.

Comparison of evaluation metrics:

Average subjectivity accuracy was 1.35% higher than its corresponding unigram feature set for unprocessed tokens. Average subjectivity accuracy was 0.45% higher than its corresponding unigram feature set for pre-processed tokens. In fact, Logistic regression produced best accuracy using subjectivity feature set on pre-processed version.

Based on NLTK classifiers and Sci-kit learn algorithms outputs, highest precision, recall and f1-measures was obtained for Sci-kit learn algorithms. The reason for this might be because the logic for Naïve Bayes is based on naïve algorithm wherein each of the feature is take multiple times i.e., the train data has some reoccurrences of data because of which redundant learning is made by the Naïve Bayes algorithm however, this is not the case with Logistic regression and Decision tree classifier algorithms. Therefore, in our case Sci-kit learn algorithms work best.

Conclusion:

The maximum accuracy attained is 70%, we obtained this accuracy for combined feature sets where we ran the Naïve Bayes classifier. After running through the entire analysis, we concluded that grouping the target variables into three categories such as Negative, Neutral and Positive would give more accuracy, however, changing the target variable composition is something which we need to consider in extreme situations.

Lessons Learned:

- Learned how to combine and implement most of the concepts learned during the semester.
- Leveraged python file system to write reusable code. Gained experience in dealing with problems that arise from working with large datasets.
- Observed how combining various feature sets affects the accuracy of the model.

Learned how different machine learning models are implemented and how their outputs differ, and what causes these differences in output.